

Assignment 3: All Things Multiprocessing

We've progressed through a good amount of material with multiprocessing, pipes, and interprocess communication, and by Wednesday we'll learn about signals and signal handlers.

Rather than building one large program, I'd like you to code up a few different things with the idea that you'll learn more by tackling multiple problems and leveraging your understanding of the material in multiple domains. All of these programs should be coded directly within a single repository, which you can get by typing the following:

```
myth02> hg clone /usr/class/cs110/repos/assign3/$USER assign3
```

There are four problems in total, and by the end of Friday's lecture, you'll be outfitted with most of the material needed to tackle the first three of them without much drama. The final problem—the prime factorization farm—will require material we won't cover until the end of Monday, but I suspect the first three will keep you busy until then. The good news is that you have 12 days to get this assignment up and operational.

Due Date: Monday, October 23, 2017 at 11:59 p.m.

Short Etude (Op. 25, No. 9): Implementing pipeline in C

Your first task is to implement the `pipeline` function. This `pipeline` function accepts two argument vectors, and assuming both vectors are legit, spawns off twin daughter processes with the added bonus that the standard output of the first is directed not to the console but rather to the standard input of the second. Here's the interface you're coding to:

```
void pipeline(char *argv1[], char *argv2[], pid_t pids[]);
```

For simplicity, you can assume that all calls to `pipeline` are well-formed and will work as expected. In other words, `argv1` and `argv2` are each valid, NULL-terminated argument vectors, and that `pids` is the base address of an array of length two. Further assume that all calls to `pipe`, `dup2`, `close`, `execvp`, and so forth succeed so that you needn't do any error checking whatsoever. `pipeline` should return without waiting for either of the child processes to finish (i.e. your `pipeline` implementation should not call `waitpid` anywhere), and the ids of the daisy-chained processes are dropped into `pids[0]` and `pids[1]`. Also, ensure that the two processes are running in parallel as much as possible, so that `pipeline({"sleep", "10", NULL}, {"sleep", "10", NULL}, pids)` takes about 10 seconds instead of 20.

You should place your implementation of `pipeline` in `pipeline.c`, and you can rely on `pipeline-test.c` and the `pipeline-test` executable it compiles to exercise your implementation.

The `pipeline-test.c` test harness I start you off with is small, so you should add many more tests of your own to prove that your `pipeline` is coded to specification.

Note that this first problem is standalone and doesn't contribute to anything else that follows (although the concept of a pipeline will come back in Assignment 4).

Short Etude (Op. 25, No. 8): Implementing subprocess in C++

Your next task is to implement an even more flexible `subprocess` than that implemented in lecture. The most important part of the `subprocess.h` interface file is right here:

```
/**
 * Function: subprocess
 * _____
 * Creates a new process running the executable identified via argv[0].
 *
 * argv: the NULL-terminated argument vector that should be passed to the
 *       new process's main function
 * supplyChildInput: true if the parent process would like to pipe
 *                   content to the new process's stdin, false otherwise
 * ingestChildOutput: true if the parent would like the child's stdout to
 *                    be pushed to the parent, false otherwise
 */
subprocess_t subprocess(char *argv[],
                        bool supplyChildInput,
                        bool ingestChildOutput) throw (SubprocessException);
static const int
kNotInUse = -1;

struct subprocess_t {
    pid_t pid;
    int supplyfd;
    int ingestfd;
};

/**
 * Function: subprocess
```

```

* -----
* Creates a new process running the executable identified via argv[0].
*
*   argv: the NULL-terminated argument vector that should be passed to the
*         new process's main function
*   supplyChildInput: true if the parent process would like to pipe
*                     content to the new process's stdin, false otherwise
*   ingestChildOutput: true if the parent would like the child's stdout to
*                      be pushed to the parent, false otherwise
*/
subprocess_t subprocess(char *argv[],
                        bool supplyChildInput,
                        bool ingestChildOutput) throw (SubprocessException);

```

Read through the `subprocess.h` documentation to see how this new `subprocess` should work, and place your implementation in `subprocess.cc`. Should any of the system calls needed to implement your `subprocess` routine fail (either in the parent or in the child), you should throw a `SubprocessException` around an actionable error message. Inspect the `subprocess-exceptions.h` file for the full, inlined definition.

Use the test harness supplied by `subprocess-test.cc` (the `.cc` extension means that that code within it is C++) to exercise your implementation, and by all means add to the `subprocess-test.cc` file to ensure that your implementation is bulletproof. When looking at `subprocess-test.cc`, you'll also get a little bit of a reminder how `try/catch` blocks work. Be sure to add your own tests to `subprocess-test.cc` to ensure that all the `(true, true)`, `(true, false)`, `(false, true)`, and `(false, false)` combinations for `(supplyChildInput, ingestChildOutput)` all work as expected.

Note that your implementation here is formally C++, since the two larger exercises that follow this one are also to be written in C++, and they each need to link against your `subprocess` implementation without drama. We're switching to C++ pretty much from this problem forward, because C++ provides better support for strings and generics than C does. C++ also provided native support for some threading and concurrency directives we'll be relying on a few weeks, and I'd rather ease you into the language now than do so when we branch into the multithreading topic. Truth be told, your C++ implementation of `subprocess` will look as it would have in pure C, save for the fact that you're throwing C++ exceptions to identify errors.

Your fully functional `subprocess` routine is used by code I wrote for the next exercise (the one requiring you to implement `trace`) and by the starter code I've given you for the final exercise of the entire assignment (the one requiring you implement the prime factorization `farm`).

Long Etude (Op. 10, No. 4): Implementing `trace` in C++

`trace` is a systems programming tool that helps us profile the execution of a secondary process and present information about all of the system calls—that is, function calls into the kernel—that the secondary executable makes. Specifically, the secondary process makes calls to system calls you certainly know about (e.g. `open`, `stat`, `read`, `write`, `close`, `sleep`) and ones you probably don't (`mmap`, `mprotect`, `ioctl`, `recv`, `getdents`).

The process running `trace` is called the **tracer**, and the process being profiled is called the **tracee**.

`trace` can be invoked in two different modes: simple and full (and your implementation needs to support both). When run in simple mode, `trace` publishes information about all of the tracee's system calls via a bare-bones presentation. Only system call codes and raw return values are posted—nothing about arguments, data types, or specific error information. To illustrate, consider the following program (you'll see it in your repo as `simple-test5.cc`):

```
int main(int argc, char *argv[]) {
    write(STDOUT_FILENO, "12345\n", 6);

    int fd = open(__FILE__, O_RDONLY);
    write(fd, "12345\n", 5);
    close(fd);
    read(fd, NULL, 64);
    close(/* bogusfd = */ 1000);
    return 0;
}
```

Assuming the above has been compiled into an executable called `simple-test5`, its bare-bones trace might look like this:

```
myth5> ./trace --simple ./simple-test5
syscall(59) = 0
syscall(12) = 35303424
// many lines omitted for brevity
syscall(1) = 12345
6
syscall(1) = 6
syscall(2) = 3
syscall(1) = -9
```

```
syscall(3) = 0
syscall(0) = -9
syscall(3) = -9
syscall(231) = <no return>
Program exited normally with status 0
myth5>
```

There are a lot of magic numbers there, but I promise that the numbers in parentheses are system call numbers (59 is for `execve`, 0 is for `read`, 1 is for `write`, 2 is for `open`, 3 is for `close`, 12 is for `brk`, 231 is for `exit_group`) and the numbers after the equals signs are return values (that 6 is the number of characters published by `write`, the -9's express `write`'s, `read`'s, and `close`'s inability to function when handed closed, incompatible, or otherwise bogus file descriptors, and `exit_group` never returns (gulp!)).

When run in full mode (i.e. without the `--simple` flag), `trace` pulls out all of the stops and prints oodles of information about each of calls:

```
myth5> ./trace ./simple-test5
execve("./simple-test5", 0x7ffeb3d1a460, 0x7ffeb3d1a470) = 0
brk(NULL) = 0xbbcffe04
// many lines omitted for brevity
write(1, "12345
", 6) = 12345
6
open("simple-test5.cc", 0, 6) = 3
write(3, "12345
", 5) = -1 EBADF (Bad file descriptor)
close(3) = 0
read(3, NULL, 64) = -1 EBADF (Bad file descriptor)
close(1000) = -1 EBADF (Bad file descriptor)
exit_group(0) = <no return>
Program exited normally with status 0
myth5>
```

You can see the return values, if negative, are always -1, but that some `errno` value is printed after that in `#define` constant form, followed by a specific error message. It turns out that `#define` constant is always mapped to the absolute value of the system call's return value (so in this case, 2), and you can easily get the more detailed error message by passing the constant (in this case 2) to the `strerror` function.

Understanding ptrace

A key function needed for this groovy piece of software to work is called `ptrace`, and you can get an abundance of information about it by typing `man ptrace` at the command line. The `ptrace` function is itself a system call, and our `trace` tool relies on it to monitor and even manipulate the execution of the tracee, inspect the contents of the tracee's virtual address space, and even inspect the contents of the registers (e.g. `%rax`, `%rsi`, `%rsp`, and so forth).

The full prototype of `ptrace` looks like this:

```
long ptrace(enum __ptrace_request request, pid_t pid, void *addr, void *data);
```

The first argument to `ptrace` is always some enumerated constant (e.g. `PTRACE_SYSCALL`, `PTRACE_PEEKUSER`, `PTRACE_TRACEME`), and the choice of constant dictates how many additional arguments are needed and what happens as a side effect of calling it. You can read through the `man` page for a description of all of the different constants, but I summarize below the list of constants you'll need for a working solution.

- `PTRACE_TRACEME`: Used by the tracee to state its willingness to be manipulated and inspected by its parent. No additional arguments are required, so a simple call to `ptrace(PTRACE_TRACEME)` does the trick.
- `PTRACE_PEEKUSER`: Used by the tracer to inspect and extract the contents of a tracee's register at the time of the call. Only the first three arguments are needed, and any value passed through `data` is ignored. A call to `ptrace(PTRACE_PEEKUSER, pid, RSI * sizeof(long))`, for instance, returns the contents of the tracee's `%rsi` register at the time of the call (provided the supplied `pid` is the process id of the tracee, of course). There are constants for all registers (`RAX`, `RSI`, `RBX`, `RSP`, etc), and the third argument is supposed to be scaled by the size of a word on that processor (which is, by definition, the size of a `long`). If the tracee temporarily freezes at the time a system call is made, then the registers contain information about what the system call is (because a number representing the system call is passed through via `%rax`) and what its arguments are (because the integers, strings, and pointers passed through as arguments sit in `%rdi`, `%rsi`, `%rdx`, `%r10`, `%r8`, and `%r9`). Of course, some system calls don't take any arguments (e.g. `fork`), some take just one (e.g. `sleep`, `close`), and some require more (`dup2` requires 2; `read`, `write`; `waitpid` require three; and a few like `pselect` require all six). When a system call returns, its return value is either an integer or a pointer, and those contents reside in `%rax`.
- `PTRACE_PEEKDATA`: Used by the tracer to inspect and extract the word of data residing at the specified location within the tracee's virtual address space. A call to `ptrace(PTRACE_PEEKDATA, pid, 0x7fa59a8b0000)` would return the eight bytes residing at address `0x7fa59a8b0000` within the tracee's virtual address space, and a call to `ptrace(PTRACE_PEEKDATA, pid, ptrace(PTRACE_PEEKUSER, pid, RDI * sizeof(long))` would return the eight bytes residing at another address, which itself resides in `%rdi`. If you know the contents of a register is an address interpretable as the base address of a `'\0'`-terminated C string, you can collect all of the characters of that string by a sequence of `PTRACE_PEEKDATA` calls.
- `ptrace(PTRACE_SYSCALL, pid, 0, 0)` instructs the tracee to continue with the understanding that the tracee will halt as it enters a system call (i.e. when the registers have been pop-

ulated with a system call code and all of its arguments) or returns from one (i.e. when the system call's return value has been dropped into `%rax`. And because the tracee halts, it's easy for the tracer to halt via a strategic call to `waitpid` until it's time for the register set to be inspected. From the tracer's point of view, the tracee will be stopped by a `SIGTRAP`, which is the signal used as part of the system function call and return to transition between user mode and kernel mode).

- `ptrace(PTRACE_SETOPTIONS, pid, 0, PTRACE_O_TRACESYSGOOD)` instructs the kernel to set bit 7 of the wait status to be 1 for all `SIGTRAP`s sent to the tracer because a system call was made. A `SIGTRAP` signal can be sent for many reasons (e.g. the tracee calls `raise(SIGTRAP)` as part of execution, or the tracee might be some gdb-like process that itself uses `ptrace` and relies on `SIGTRAP`s to be fired at it because a value being watched changed, or because execution hit some set breakpoint). Bottom line: `PTRACE_SETOPTIONS` and `PTRACE_O_TRACESYSGOOD` help us distinguish `SIGTRAP`s generated by system calls from other `SIGTRAP`s.

Here is a key part of the `ptrace` man page (slightly edited for clarity):

A process can initiate a trace by calling `fork` and having the resulting child do a `PTRACE_TRACEME`, followed by a `raise(SIGSTOP)` (which causes the child process to self-halt until instructed by the tracer to continue), followed by the `execvp`. While being traced, the tracee will stop each time a signal is delivered, even if the signal is being ignored. The tracer will be notified at its next call to `waitpid`; that call will return a status value containing information that indicates the cause of the stop in the tracee. While the tracee is stopped, the tracer can use various `ptrace` requests to inspect the tracee. The tracer then prompts the tracee to continue.

Tips and Tidbits

- When run in `full` mode, you'll need access to a collection of maps storing system call function names, system call signatures, and return types. You'll also need to know what `#define` constants should be printed (in text form) when system calls fail. I'm giving you a bunch of libraries that crawl over a subset of the system header files to extract information about system call numbers and `errno` constants, and over a reasonably large number of source files implementing the Linux kernel. You're welcome to peruse the `trace-error-constants.cc` and `trace-system-calls.cc` source files, but you really only need to read through the corresponding interface files to understand what they do for you. (The function that crawls over the Linux kernel code base in `/usr/src/linux-source-3.13.0/linux-source-3.13.0` assumes a working `subprocess` implementation, so you'll need to make sure you get that working before you can expect the support libraries to work).
- The very first time you run `trace`, you should expect it to take a very long time to read in all of the prototype information for the linux kernel source tree. All of the prototype information is cached in a local file after that (the cache file will sit in your repo with the name `.trace_signatures.txt`), so `trace` will fire up much more quickly after that. Should you want to rebuild the prototype cache for any reason, you can invoke `trace` with the `--rebuild` flag, as with `trace --rebuild make clean`.
- For the purposes of this assignment, you can assume that all return values should be printed as integers, except for `brk` and `mmap`, which we'll assume return `void *`'s. This information

isn't as easily extracted from the headers or kernel source as you might think, so I'm going with a massive simplification here.

- There are a collection of string utility functions I wrote that reside in `/usr/class/cs110/local/include/string-utils.h`, and you can use them by `#include'ing string-utils.h`.
- The x86_64 architecture requires that system call function codes and return values be placed in `%rax`. System call arguments are placed in `%rdi`, `%rsi`, `%rdx`, `%r10`, `%r8`, and `%r9` as needed. To be clear, the first argument if needed is placed in `%rdi`; the second argument, if needed, is placed in `%rsi`; and so forth.
- Earlier, I mentioned the system call function code is passed in via `%rax` and that's true. But by the time that necessary `PTRACE_PEEKUSER` call gets made, the true `%rax` has been clobbered. Fortunately, the solution is to, for just this one scenario, rely on `ORIG_RAX` instead, which still houses the system call code at the time you need it.
- One particularly tricky part: printing a system call argument when that argument is a string. The base address of that string's leading character will sit in a register, and you'll need to use a combination of `PTRACE_PEEKUSER` and `PTRACE_PEEKDATA` to reconstitute the string so that it can be printed.
- All arguments are either ints, `char *`s, or general pointers. The `long` returned by `ptrace` needs to be truncated to an `int`, converted to a C++ `string`, or reinterpreted as a `void *` before printing it.
- The `sys/reg.h` header defines a collection of constants used to represent registers, and they are always all-caps versions of the registers they represent: `RAX`, `RSI`, `RBX`, etc.
- This particular program certainly relies on signals, but there are no exposed signal handler functions in my own solution. We won't get to signals and signal handlers until Wednesday, but that shouldn't block you from starting up on this assignment.
- The return value of `trace` is always the return value of the tracee. You needn't worry about tracees that exit because of some uncaught signal, as that requires more advanced `ptrace` work, because tracees don't always properly report their own death via `WIFEXITED` and `WIFSIGNALED` when terminated by a signal.
- My own solution has been compiled into an executable named `trace_soln`, and it's sitting in `/usr/class/cs110/samples/assign3`. Feel free to play with it and compare what it generates to what yours does. Your `assign3` repo contains a symlink called `slink`, which is an alias for `/usr/class/cs110/samples/assign3`.

Scherzo No. 4: Implementing farm in C++

Your final challenge is to harness the power of a computer's multiple cores to manage a collection of executables, each running in parallel to contribute its share to a larger result. For the purposes of this problem, we're going to contrive a scenario where the computation of interest—the prime factorization of arbitrarily large numbers—is complex enough that some factorizations take multiple seconds or even minutes to compute. The factorization algorithm itself isn't the focus here, save for the fact that it's potentially time consuming, and that should we need to compute multiple prime factorizations, we should leverage the computing resources of our trusty Unix cluster

machines to multiprocessing and generate output more quickly.

Consider the following Python program called `factor.py`:

```
self_halting = len(sys.argv) > 1 and sys.argv[1] == '--self-halting'
pid = os.getpid()
while True:
    if self_halting: os.kill(pid, signal.SIGSTOP)
    try: num = int(raw_input()) # raw_input blocks, eventually returns a single
line from stdin
    except EOFError: break; # raw_input throws an EOFError when EOF is detected
    start = time.time()
    response = factorization(num)
    stop = time.time()
    print ' %s [pid: %d, time: %g seconds]' % (response, pid, stop - start)
```

You really don't need to know Python to understand how it works, because every line of this particular program has a clear C or C++ analog. The primary things I'll point out are:

- Python's `print` operates just like C's `printf` (and it's even process-safe)
- `raw_input` reads and returns a single line of text from standard input, blocking indefinitely until a line is supplied
(chomping the `'\n'`) or until end-of-file is detected
- `factorization` is something I wrote; it takes an integer (e.g. 12345678) and returns the prime factorization (e.g. $12345678 = 2 * 3 * 3 * 47 * 14593$) as a string. You'll see it when you open up `factor.py` in your favorite text editor
- The `os.kill` line prompts the script to stop itself (but only if the script is invoked with the `'--self-halting'` flag) and wait for it to be restarted via `SIGCONT`

The following should convince you our script does what you'd expect (I'm using `bash` here, where the `time` builtin clocks the entire pipeline):

```
real  0m16.806s
user  0m16.793s
sys   0m0.008smyth02> printf "1234567\n12345678\n" | ./factor.py
1234567 = 127 * 9721 [pid: 14391, time: 0.100041 seconds]
```

```

12345678 = 2 * 3 * 3 * 47 * 14593 [pid: 14391, time: 1.03848 seconds]
myth02> time printf "1234567\n12345678\n123456789\n1234567890\n" | ./factor.py
1234567 = 127 * 9721 [pid: 14440, time: 0.108153 seconds]
12345678 = 2 * 3 * 3 * 47 * 14593 [pid: 14440, time: 1.04659 seconds]
123456789 = 3 * 3 * 3607 * 3803 [pid: 14440, time: 10.1917 seconds]
1234567890 = 2 * 3 * 3 * 5 * 3607 * 3803 [pid: 14440, time: 102.537 seconds]

real    1m53.911s
user    1m53.929s
sys      0m0.033s

myth02> printf "1001\n10001\n" | ./factor.py --self-halting
myth02> kill -CONT %1
1001 = 7 * 11 * 13 [pid: 15973, time: 0.000144005 seconds]
myth02> kill -CONT %1
10001 = 73 * 137 [pid: 15973, time: 0.000889063 seconds]
myth02> kill -CONT %1
myth02> kill -CONT %1
kill: No such job.

myth02> time printf "123456789\n123456789\n" | ./factor.py
123456789 = 3 * 3 * 3607 * 3803 [pid: 2143, time: 8.39598 seconds]
123456789 = 3 * 3 * 3607 * 3803 [pid: 2143, time: 8.39575 seconds]

real    0m16.806s
user    0m16.793s
sys      0m0.008s

```

This last test may look silly, but it certainly verifies that one process is performing the same factorization twice, in sequence, so that the overall running time is roughly twice the time it takes to compute the factorization the first time (no caching here, so the second factorization does it all over again).

My `factorization` function runs in $O(n)$ time, so it's very slow for some large inputs. Should you need to compute the prime factorizations of many large numbers, the `factor.py` script would get the job done, but it may take a while. If, however, you're `ssh`'ed into a machine that has multiple processors and/or multiple cores (the `myths` have two cores each, and the `ryes` have eight!), you can write a program that manages several processes running `factor.py` and tracks which processes are idle and which processes are deep in thoughtful number theory.

You're going to write a program—a C++ program called `farm`—that can run on the `myths` or the `ryes`. `farm` will spawn several workers—one for each core, each running a self-halting instance of `factor.py`, read an unbounded number of positive integers (one per line, no error checking required of you for this problem either), forward each integer on to an idle worker (blocking until one or more workers is ready to read the number), and allow all of the workers to cooperatively publish all prime factorizations to standard output (without worrying about the order in which they're printed). To illustrate how `farm` should work, check out the following test case:

```
real 0m10.667s
user 0m41.197s
sys 0m0.099s
rye01> time printf "123456789\n123456789\n123456789\n123456789\n"
| ./farm

There are this many CPUs: 8, numbered 0 through 7.
Worker 25528 is set to run on CPU 0.
Worker 25529 is set to run on CPU 1.
Worker 25530 is set to run on CPU 2.
Worker 25531 is set to run on CPU 3.
Worker 25532 is set to run on CPU 4.
Worker 25533 is set to run on CPU 5.
Worker 25534 is set to run on CPU 6.
Worker 25535 is set to run on CPU 7.

123456789 = 3 * 3 * 3607 * 3803 [pid: 25528, time: 10.2493 seconds]
123456789 = 3 * 3 * 3607 * 3803 [pid: 25531, time: 10.3282 seconds]
123456789 = 3 * 3 * 3607 * 3803 [pid: 25530, time: 10.4229 seconds]
123456789 = 3 * 3 * 3607 * 3803 [pid: 25529, time: 10.6165 seconds]

real 0m10.667s
user 0m41.197s
sys 0m0.099s
```

Note that each of four processes took about the same amount of time to compute the identical prime factorizations, but because each of the four processes was assigned to different cores, the real (aka perceived) user time was under 11 seconds. Note that prime factorizations aren't required to be published in order, and repeat requests for the same prime factorization are all computed from scratch, without any caching.

Your `farm.cc` implementation will make use of the following C++ record, global constants, and

global variables:

```
static const size_t kNumCPUs = sysconf(_SC_NPROCESSORS_ONLN);
static vector<worker> workers(kNumCPUs); // space for kNumCPUs, zero-arg constructed workers
static size_t numWorkersAvailable = 0; struct worker {
    worker() {}
    worker(char *argv[]) : sp(subprocess(argv, true, false)), available(false) {}
    subprocess_t sp;
    bool available;
};

static const size_t kNumCPUs = sysconf(_SC_NPROCESSORS_ONLN);
static vector<worker> workers(kNumCPUs); // space for kNumCPUs, zero-arg constructed workers
static size_t numWorkersAvailable = 0;
```

The `main` function we give you includes stubs for all of the helper functions that decompose it, and that `main` function looks like this:

```
int main(int argc, char *argv[]) {
    signal(SIGCHLD, markWorkersAsAvailable);
    spawnAllWorkers();
    broadcastNumbersToWorkers();
    waitForAllWorkers();
    closeAllWorkers();
    return 0;
}
```

This final problem can be tricky, but it's perfectly manageable provided you follow this road map:

- Advance on to `spawnAllWorkers`, which spawns a self-halting `factor.py` process for each core and updates the global `workers` vector so that each worker contains the relevant `subprocess_t` allowing `farm.cc` to monitor it and pipe prime numbers to it. You can assign a process to always execute on a particular core by leveraging functionality outlined in the `CPU_SET` and `sched_setaffinity` man pages (i.e. type in `man CPU_SET` to learn about the `cpu_set_t` type, the `CPU_ZERO` and `CPU_SET` macros, and the `sched_setaffinity` function).

- Implement the `markWorkersAsAvailable` handler, which gets invoked whenever one of the child processes self-halts (prompting the kernel to SIGCHLD signal the parent). Call `waitpid` to surface the pid of the child that recently self-halted, and mark it as available.
- Implement a `getAvailableWorker` helper function, which you'll use to decompose the `broadcastNumbersToWorkers` function in the next step. You should never busy wait; instead, investigate `sigsuspend` (by typing `man sigsuspend`) as a way of blocking indefinitely until at least one worker is known to be available.
- Flesh out the implementation of `broadcastNumbersToWorkers`. I'm giving you a tiny hint here—that `broadcastNumbersToWorkers` keeps on looping until either EOF is detected (or until the `farm` user messes up and deviated from the required input format). Investigate the SIGCONT signal as the means to restart another stopped process.

```
static void broadcastNumbersToWorkers() {
    while (true) {
        string line;
        getline(cin, line);
        if (cin.fail()) break;
        size_t endpos;
        /* long long num = */ stoll(line, &endpos);
        if (endpos != line.size()) break;
        // you shouldn't need all that many lines of additional code
    }
}
```

- Implement `waitForAllWorkers`, which does more or less what it says—it waits for all workers to self-halt and become available.
- Last but not least, implement the `closeAllWorkers` routine to uninstall the SIGCHLD handler and restore the default (investigate the SIG_DFL constant), cajole all child processes to exit by sending them EOFs, and then wait for them to all actually exit.

Your implementation should be exception-safe, and nothing you write should orphan any memory. You should `ssh` into the rye machines (either `ssh rye01.stanford.edu` or `ssh rye02.stanford.edu` to exercise your implementation over there to confirm the extra parallelization that comes with eight cores over just two).

Finally, my own solution has been compiled into an executable named `farm_soln`, and it's sitting aside `trace_soln` in `/usr/class/cs110/samples/assign3`. That's another sample executable for you to play with.